

PROYECTO DEEP LEARNING



Ruben Tenreiro
Jesus Moncho
Marco Hernani
Helen Navarro



01

Carga de
imágenes

04

Entrenamiento

02

Etiquetado de
imágenes

05

Incidencias

03

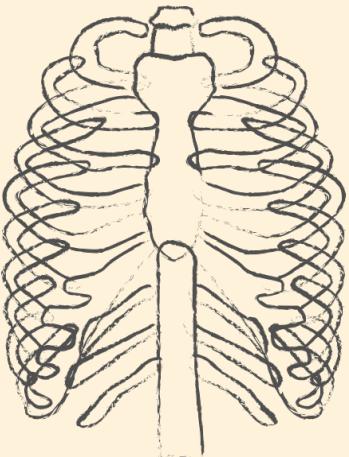
Balanceo de
clases

06

Conclusiones



INTRODUCCIÓN



Squad ETL

Nos encargaron el proyecto de detectar mediante reconocimiento de imágenes si un paciente tiene neumonía o no, con un proceso de entrenamiento de DEEP LEARNING.

Importación de librerías

```
import tensorflow as tf
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import keras
from keras.models import Sequential
from keras.layers import Dense, Conv2D , MaxPool2D , Flatten , Dropout , BatchNormalization
from keras.preprocessing.image import ImageDataGenerator
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report,confusion_matrix
from keras.callbacks import ReduceLROnPlateau
import cv2
import os
```



TensorFlow





CARGA DE IMÁGENES



MATT GROENING

**SIGUIENTE
POR FAVOR**





Cargamos las imágenes de radiografías con cv2



Las imágenes se convierten en matrices de 300×300

image_size = (300, 300)

```
1 # What these images really are
2
3 train_data[0][0]
array([[18, 19, 24, ..., 5, 2, 0],
       [19, 21, 24, ..., 6, 3, 0],
       [18, 22, 24, ..., 7, 5, 0],
       ...,
       [10, 10, 10, ..., 15, 15, 15],
       [10, 10, 10, ..., 15, 15, 15],
       [10, 10, 10, ..., 15, 15, 15]], dtype=uint8)
```

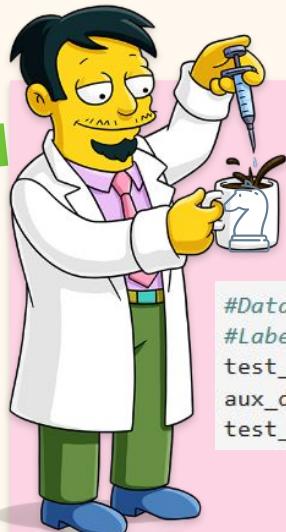


Etiquetado de ímágenes

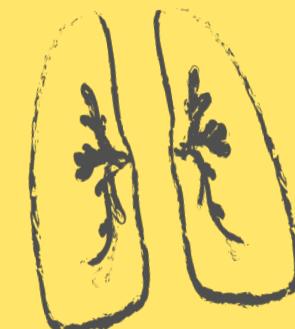


MATT GROENING

ETIQUETADO DE IMÁGENES



```
#Data needed to train our model  
#Labels, 0:normal and 1:pneumonia  
test_data = load_images("datasets/test/PNEUMONIA", image_size, 1)  
aux_data = load_images("datasets/test/NORMAL", image_size, 0)  
test_data.extend(aux_data)
```





Balanceo de clases



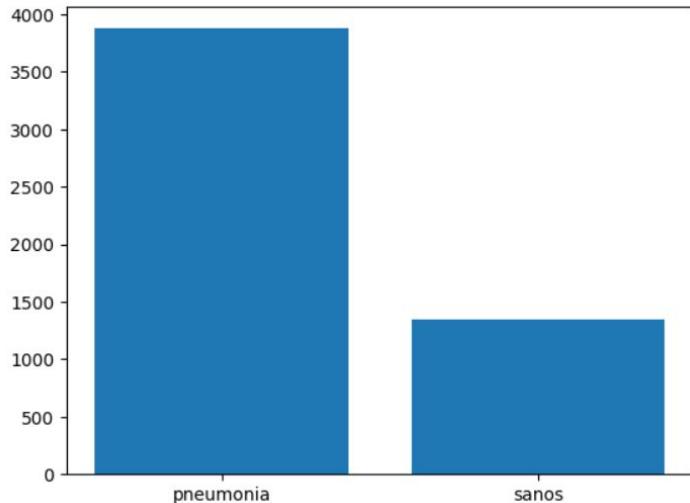
MATT GROENING

**SIGUIENTE
POR FAVOR**



Datos sin balancear

```
1 bar_data = [len(train_pneu), len(train_norm)]  
2 plt.bar(["pneumonia", "sanos"], bar_data)  
3 plt.show()
```

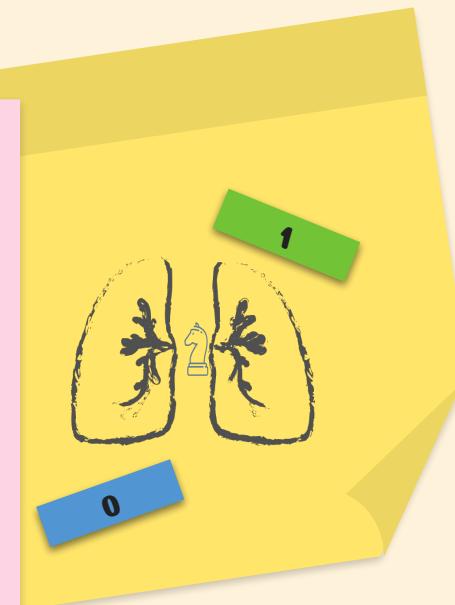
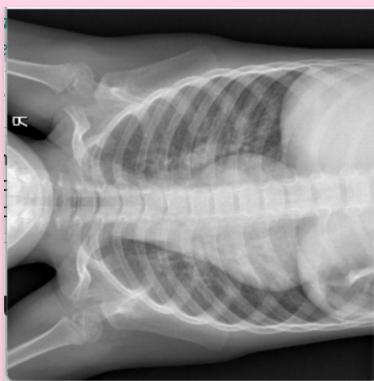


```
1 total = healthy + pneumonia  
2 print(f"sanos: {round(healthy/total*100, 2)}%")  
3 print(f"con pneumonia: {round(pneumonia/total*100, 2)}%")
```

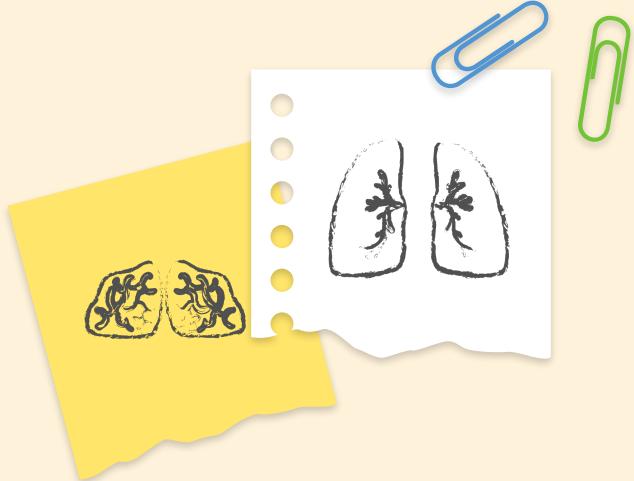
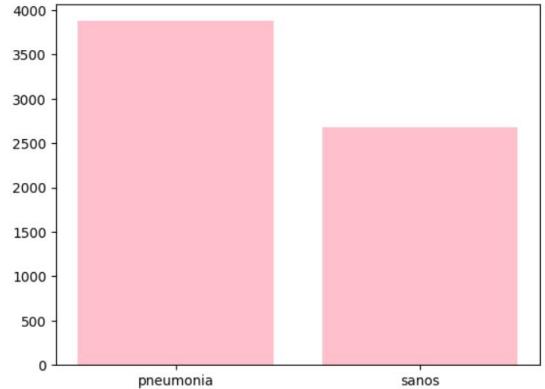
sanos: 26.97%
con pneumonia: 73.03%

Balanceo de clases

```
#We add normal data  
for x, _ in train_normal[:]:  
    train_data.append( (x.T, 0) )
```



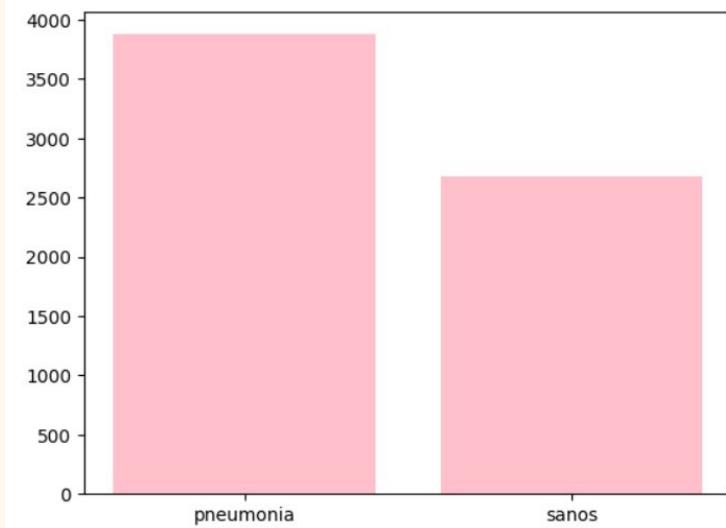
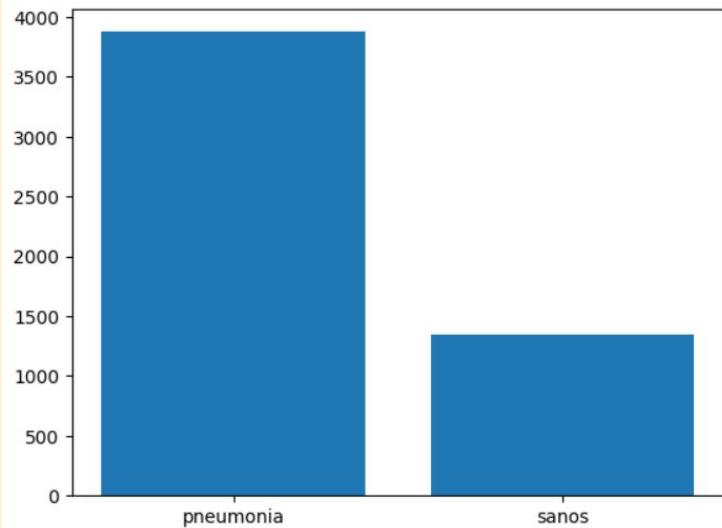
```
1 bar_data = [len(train_pneu), len(train_norm) + len(train_norm_mirror)]
2 plt.bar(["pneumonia", "sanos"], bar_data, color = "pink")
3 plt.show()
```



```
1 train = train_pneu + train_norm + train_norm_mirror
2
3 print(f"sanos: {round(len(train_norm + train_norm_mirror)/len(train)*100, 2)}%")
4 print(f"con pneumonia: {round(len(train_pneu)/len(train)*100, 2)}%")
```

sanos: 40.9%
con pneumonia: 59.1%

Comparación Inicio-final





Entrenamiento



Normalización y creación de train y test sets



```
1 #Outputs
2 train_y = []
3 #Inputs
4 train_x = []
5
6 for x,y in train_data:
7     train_y.append(y)
8     train_x.append(x)
9
10 train_y = np.array(train_y, dtype='float32')
11
12 train_x = np.array(train_x, dtype='float32')/255
```

```
1 #Outputs
2 test_y = []
3 #Inputs
4 test_x = []
5
6 for x,y in test_data:
7     test_y.append(y)
8     test_x.append(x)
9
10 test_y = np.array(test_y, dtype='float32')
11
12 test_x = np.array(test_x, dtype='float32')/255
```



Creación del modelo

```
1 model = Sequential([
2     Flatten(input_shape=image_size),
3     Dense(32, activation='relu'),
4     Dropout(0.2),
5     Dense(16, activation='relu'),
6     Dropout(0.2),
7     Dense(8, activation='relu'),
8     Dropout(0.2),
9     Dense(1, activation="sigmoid") #'softmax')
10])
11
12
13 model.compile(optimizer='adam',
14                 loss='binary_crossentropy',
15                 #loss='sparse_categorical_crossentropy',
16                 metrics=['accuracy'])
```

Entrenamiento

```
model.fit(train_x, train_y, epochs=5)
```

```
Epoch 1/5
205/205 [=====] - 6s 23ms/step - loss: 1.0538 - accuracy: 0.7423
Epoch 2/5
205/205 [=====] - 5s 22ms/step - loss: 0.5409 - accuracy: 0.7837
Epoch 3/5
205/205 [=====] - 4s 19ms/step - loss: 0.5064 - accuracy: 0.7927
Epoch 4/5
205/205 [=====] - 4s 21ms/step - loss: 0.4822 - accuracy: 0.8062
Epoch 5/5
205/205 [=====] - 4s 20ms/step - loss: 0.4466 - accuracy: 0.8280
<keras.callbacks.History at 0x1c4f97a5b10>
```

Resultados



```
model.evaluate(train_x, train_y, verbose=2)
```

```
205/205 - 1s - loss: 0.3243 - accuracy: 0.9395 - 1s/epoch - 6ms/step  
[0.3242920935153961, 0.9394540190696716]
```

```
model.evaluate(test_x, test_y, verbose=2)
```

```
20/20 - 0s - loss: 0.4868 - accuracy: 0.8205 - 150ms/epoch - 7ms/step  
[0.4867691695690155, 0.8205128312110901]
```



Incidencias



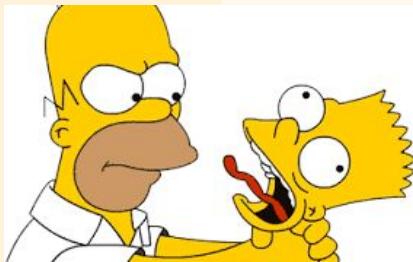
Una vez obtenido el modelo este hacia overfitting.

- Posible causa: falta de datos de personas sanas (1/4 respecto al total)
- Posible solución: crear nuevas imágenes rotando las anteriores

1/4 sanas -> aprox 2/4 2000 -> 2/4 4000

3/4 pneu -> aprox 2/4 2000-> 2/4 4000

accuracy: **~70 test, ~90 train**



Gracias por
vuestra atención.
¿Dudas?

